

Machine Learning Nanodegree

Capstone Proposal

Sam Mottahedi

01/08/2018

1 Definition

1.1 Project Overview

Natural language processing (NLP) is one of the most important technologies of the information age. Understanding complex language utterances is also a crucial part of artificial intelligence. Applications of NLP are everywhere because people communicate most everything in language: web search, advertisement, emails, customer service, language translation, radiology reports, etc. There are a large variety of underlying tasks and machine learning models behind NLP applications.

1.2 Problem Statement

Free expression and sharing information is the greatest impact of internet in modern society. Unfortunately, online abuse and harassment can lead limit self expression on the web. Many platforms struggle to effectively facilitate conversations, leading many communities to limit or completely shut down user comments.

In this [Kaggle competition](#), participants are challenged to build a multi-headed model that's capable of detecting different types of toxicity like threats, obscenity, insults, and identity-based hate.

1.3 Metrics

Submissions are evaluated on the mean column-wise log loss. In other words, the score is the average of the log loss of each predicted column.

Since class labels are not mutually exclusive, a multi-label classification loss function is required here. Multi-label classification (MLC) is a prediction problem in which several class labels are assigned to single instances simultaneously as follows:

$$loss(\hat{y}, y) = \frac{1}{|L|} \sum_{l=1}^{l=|L|} -(y_l - \log(\hat{y}_l) + (1 - y_l) \cdot \log(1 - \hat{y}_l)) \quad (1)$$

2 Analysis

2.1 Data Exploration

The [Kaggle competition dataset](#) provided is Wikipedia Human Annotations of Toxicity on Talk Pages and contains 160,000 human labelled annotations based on asking 5000 crowd-workers to rate Wikipedia comments according to their toxicity (likely to make others leave the conversation). Each comment was rated by 10 crowd-workers. The Test dataset which is used to evaluating performance in competition consist of 226,998 unlabeled comments. Each comment in the training set can be labeled with 6 labels which are not mutually exclusive.

The toxic comment data set have the following fields:

- "id": (string)
- "comment texts": comments (string)
- "toxic": toxic comment label (binary)
- "sever toxic": severely toxic comment label (binary)
- "obscene": obscene comment label (binary)
- "threat": threatening comment label (binary)
- "insult": insulting comment label (binary)
- "identity hate": identity hate comment label (binary)

2.2 Exploratory Visualization

It is important to know how labeled comment are distributed in the dataset and wether or not they the labels as balanced. Figure 1 shows the distribution of labels

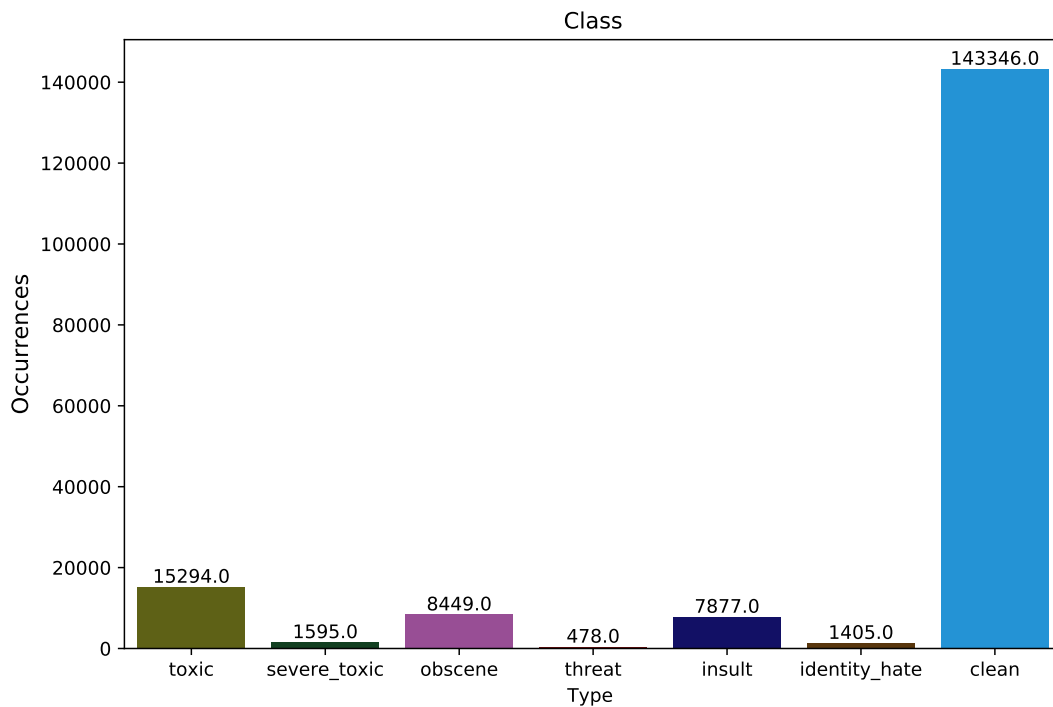


Figure 1: Toxic comment dataset label distribution

The plot shows that most of the comments are innocent and small number of comments are classified as any form toxic comment. The comments classified as severely toxic, threat and identity hate are rare compared to toxic, obscene and insult.

An example of a comment is provided bellow with was labeled as both toxic and a threat.

In this project, each comment is treated as set of features using pre-trained GloVe word embedding. A Bi-directional recurrent neural network with Gated Recurrent Units (GRU). Attention mechanism is added on top of the Bi-directional RNN which improve the performance of the model focusing on important parts of long sequence and reduce the effect of noise and unrelated information. The classification task is a multi-label classification where comments toxicity labels are not mutually exclusive.

2.4 Benchmark

The baseline chosen here in order to get better understanding of the problem from a general machine-learning perspective. To this end, the baseline model is chosen to be a logistic regression model based on Term Frequency-inverse document frequency data which can achieve column-wise log-loss of 0.05567.

3 Methodology

3.1 Data Processing

The preprocessing is mostly done using functions available in data.py module. The preprocessing step before training or testing consist of following steps:

- Preprocessing comments:
 - tokenizing
 - converting to lower-case letters
 - removing stopwords
 - normalizing numbers, date, ...
- Vocabulary: creating dictionary of words with at least $count() > Thresh_{hold}$.
- Index to word: dictionary with keys equal to word index and values being a word in the vocabulary
- Word to index: dictionary with keys equal to vocabulary's word and value being words index.
- Sequence index file: a file containing sequence of word index for each comment.

During training, the pre-processed comment are read from sequence to index file and push to TensorFlow input placeholders. In addition, at the beginning of the training the word vectors corresponding to the words in the vocabulary are extracted from pre-trained GloVe word vector with 6 Billion words and specified word vector dimension.

The preprocessing process has the following adjustable parameters:

- Word frequency threshold
- GLoVe word vector dimension (50, 100, 200)
- Using trainable or not trainable word vector tensor
- Maximum input sequence length




Figure 3: Bi-directional RNNs

3.2 Implementation

3.2.1 Text Classification with Recurrent Neural Networks

The architecture used in this work is Recurrent Neural Network (RNN) that takes the sequence of words index and an encode the information in the text in the last output layer. The last output layer of the RNN is passed through two feed forward layers and a final 1×6 layer where a element wise sigmoid function is applied and a multi-label predictions are generated. Since some of the comments are quite long, a Bi-directional RNN is used which has two Gated Recurrent Cells (GRU) for forward and backward processing of the input sequence and outputs the concatenated output oof the forward and backward cell [figure 3].

3.3 Refinement

3.3.1 Attention Mechanism

Since not all words contribute equally to the representation of the sequence meaning we need to use an attention mechanism [Yang et al. (2016)] that extract important words to the meaning of the sequence and aggregates the presentation of informative words to form a sentence vector from [figure 4]

$$u_{it} = \tanh(W_w h_{it} + b_w) \quad (2)$$

$$\alpha_{it} = \frac{\exp(u_{it}^T u_w)}{\sum_t \exp(u_{it}^T u_w)} \quad (3)$$

$$o = \sum_t \alpha_{it} h_{it} \quad (4)$$

where the word annotation h_{it} is passed to a one-layer MLP to get u_{it} as hidden representation h_{it} the importance of the word is measured by comparing u_{it} to a word context vector u_w and then passed through a softmax function to get normalized importance weight α_{it} and the output vector is weighted sum of word annotation based of the weights. The context vector is randomly initialized and jointly learned during the training process.

The completed TensorFlow computational graph can be seen in Figure [5]

3.3.2 Other Measures

In addition to the attention layer other measures such as dropout layers, exponential learning rate decay were used to reduce the log-loss and improve the model generalization.

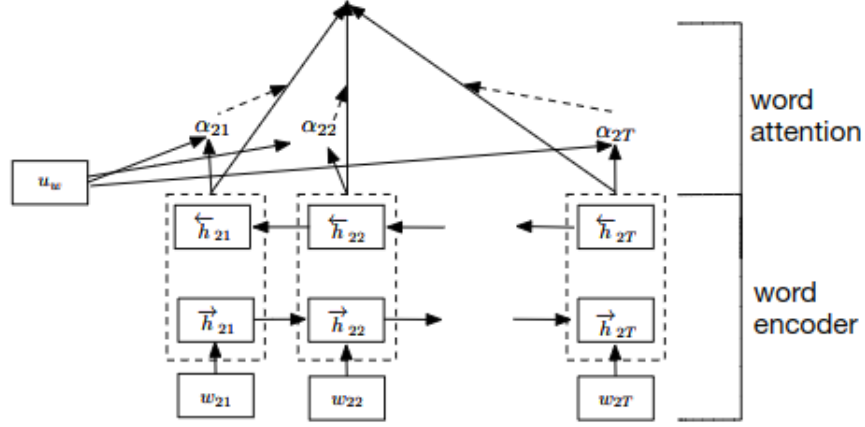


Figure 4: Attention Mechanism Yang et al. (2016)

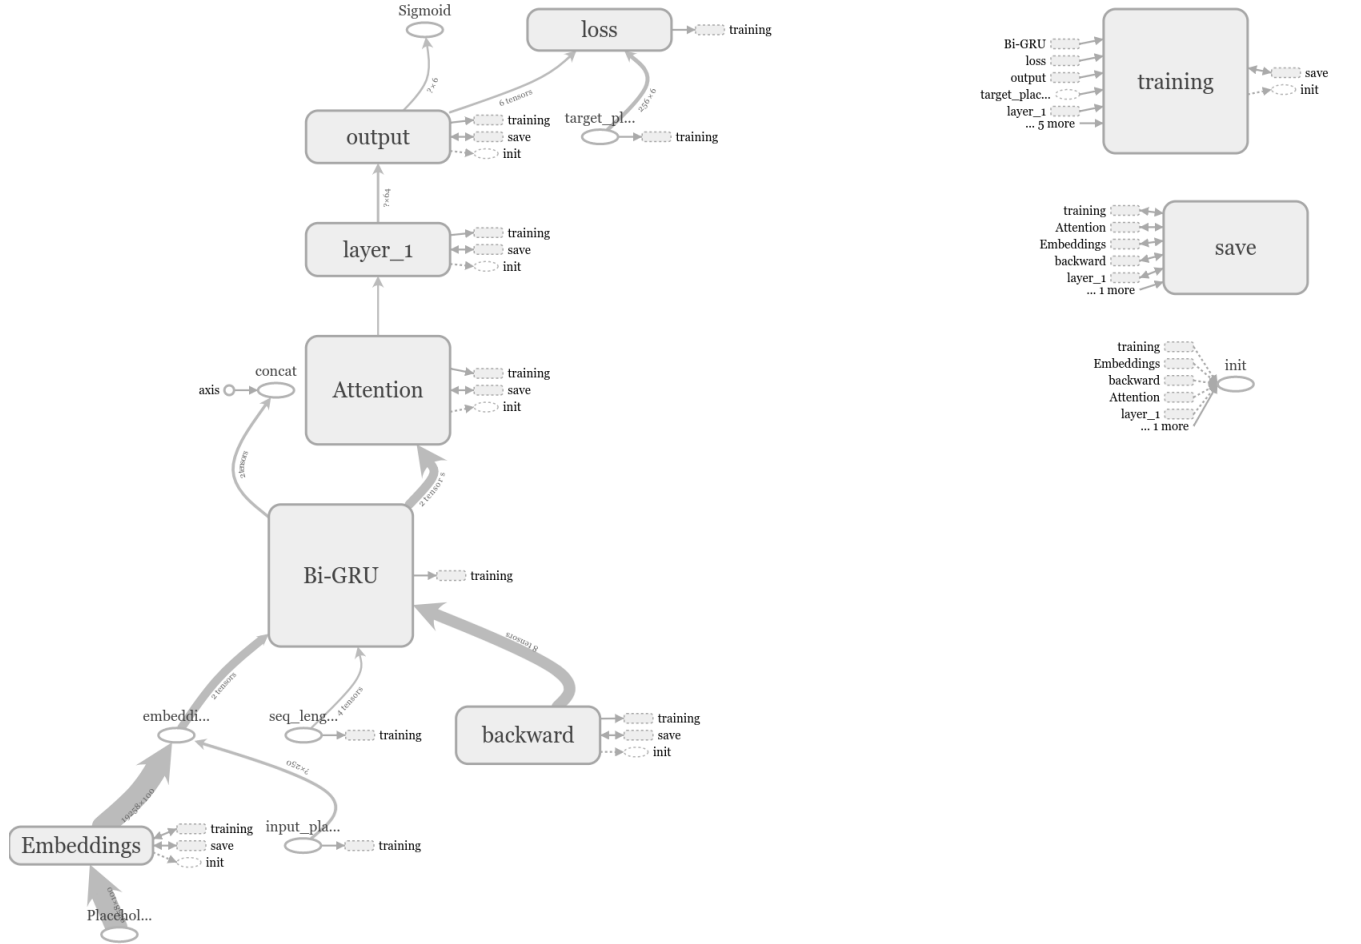


Figure 5: TensorFlow computational graph for sequence classification model with attention mechanism

4 Results

4.1 Model Evaluation and Validation

During development process, the training data was split into training (90%) and validation (10%) sets. Since this project is part of Kaggle competition, instead of using a test set the submission result is used as a metric for final evaluation of the model.

The final description of the final model is as follows:

- Batch Size = 256
- Embedding dimension = 100
- Number of GRU Cells = 128
- Feed-forward layers = 128
- Attention size = 256
- Maximum input sequence length = 250

An early-stopping mechanism is implemented to automatically save model weights corresponding to best evaluation log-loss and stop the training to prevent over-fitting to the training set.

4.2 Justification

5 Conclusion

5.1 Free-Form Visualization

5.2 Reflection

5.3 Improvement

References

Zichao Yang, Diyi Yang, Chris Dyer, Xiaodong He, Alex Smola, and Eduard Hovy. Hierarchical attention networks for document classification. In *Proceedings of the 2016 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 1480–1489, 2016.